

Foundations of All Concepts Implemented in This Project

Nicolas Becker

September 6, 2026

Contents

1	Abstract	3
2	Introduction	3
2.1	GAN Issues	3
3	Architectures	3
3.1	Vanilla GAN	3
3.2	DCGAN	4
3.3	Resnet-GAN	4
3.4	Conditional GAN	4
3.5	Super Resolution GAN	4
3.6	BigGAN	4
3.7	Summary and Comparison	4
4	Objective Functions	4
4.1	Least Square Loss	4
4.2	Wasserstein GAN	4
4.3	Wasserstein GAN with Gradient Penalty	4
4.4	Hinge GAN	4
4.5	Summary and Comparison	4
5	Regularizers	4
5.1	Feature Matching	4
5.2	Minibatch-Discrimination	5
5.3	Spectral Normalization	5
5.4	Orthogonal Normalization	5
5.5	Summary and Comparison	5
6	Experiments	5
6.1	Methods	5
7	Combinations tried	5
7.1	Limitations	5
8	Training	5
9	Summary	7

1 Abstract

This project will cover several GAN

2 Introduction

In 2014 a certain type of generative models was invented by **goodfellow2014generative**, these new type of neural networks are called *Generative Adversarial Networks* (GANs). Compared to new generative models like **DALL-E** developed by Open-AI, which is based on *Diffusion Models* (These models will probably come also in a own project as soon as GANs are accomplished), their ability of creating diverse and realistic looking images is limited. This comes with the cost of relevance in modern applications. However, GANs are still used in a wide range of applications from data augmentation and self supervised learning. One of the biggest advantage still remains in the fast inference compared to the diffusion models.

2.1 GAN Issues

Training GANs is inherently complex firstly due to their non convergence nature. This is easily seen by equation 1, their objective is against each other. Therefore the training is instable and may not converge til a global minima.

Secondly, due to their ability to be stuck in a local minima and only produce limited modes of the actual training data. The generator is only producing partiell parts of the training data. For instance, imagine that the generator is only producing the number 6 and 7 but not the others. Unfortunately, if the discriminator is not well trained only some nodes are sufficiently to fool the discriminator. This is called Mode colapse.

3 Architectures

3.1 Vanilla GAN

The beginning of this hopefully reasonably overview, will be the start of GANs with **goodfellow2014generative**, which might be the most influential paper about GANs. Shortly said, GANs consist of two neural networks working against each other to surpass the performance of the counterpart. One of the participants is the Generator that tries to generate data that resamples the original data. This procedure consists on the assumption that each dataset can be seen as a probability distribution. Following, this assumption the generator tries to learn the distribution of the dataset and can sample realistic looking images from the learned distribution. The Discriminators' objective is to distinguish whether real or generated data (fake data) was fed into the discriminator. The Generator learns to fool the discriminator, while the discriminator has to learn to distinguish. Mathematically, this can be seen with expression:

$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p_{\text{data}}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log (1 - D(G(z)))] . \quad (1)$$

the objective of the discriminator is actually to maximize the probability of whether beeing real in the first term or zero in the second term.

3.2 DCGAN

Based on the equations like (1), DCGAN is equivalent in the computation but differ only in the architecture. DCGAN stands for *Deep Convolutional GAN*, better said they are based on convolutional layers, thereby the ability to resample photo realistic images is increased. The invention of DCGAN was made by **radford2016dcgan**. Additionally, they found good hyperparameters to avoid instable training. New opportunities came across due to their convolutional nature. For instance, it became possible to use interpolation within the latent space, which makes it possible to linearly interpolate styles across a dataset. The main application from the paper invented interpolation with celebA to change semantic features of the faces.

Since convolutional layers introduce an inductive bias while processing images the quality of the generated images increases. Unfortunately, the quality is still limited and the training process is also instable.

3.3 Resnet-GAN

Training deep networks is heavy, but with the rise of resnet, the ability of training became more possible. The speciality of resnet is the skip connection between the input of the layer and the output. Implementarily, the input is added to the output, so that the model only has to learn the difference, thus the gradient flow is more fluently.

3.4 Conditional GAN

3.5 Super Resolution GAN

3.6 BigGAN

3.7 Summary and Comparison

4 Objective Functions

4.1 Least Square Loss

4.2 Wasserstein GAN

4.3 Wasserstein GAN with Gradient Penalty

4.4 Hinge GAN

4.5 Summary and Comparison

5 Regularizers

5.1 Feature Matching

The idea is to condition the discriminator in rather distinguishing between real and fake data to match the features of a intermediate layer of the discriminator for real and fake data.

5.2 Minibatch-Discrimination

5.3 Spectral Normalization

5.4 Orthogonal Normalization

5.5 Summary and Comparison

6 Experiments

6.1 Methods

7 Combinations tried

An exemplary overview of the structure that will follow soon, some trainings are in progress.

Table 1: Overview of the evaluated GAN configurations.

ID	Architecture	Loss	Regularization
A	DCGAN	Vanilla GAN	None
B	DCGAN	LSGAN	Spectral Normalization
C	DCGAN	LSGAN + Feature Matching	Spectral Normalization
D	DCGAN	LSGAN	Minibatch Discrimination
E	DCGAN	LSGAN	Label Smoothing
F	DCGAN	LSGAN	Spectral Normalization + Label Smoothing
G	DCGAN	LSGAN	Spectral Normalization + Minibatch Discrimination
H	DCGAN	WGAN	Weight Clipping
I	DCGAN	WGAN-GP	Gradient Penalty

7.1 Limitations

Unfortunately, training GANs is highly time and computation consuming. Therefore not all trainings are done until optimum. The trainings computed here are invariably done on one Nvidia RTX 4090. Trainings that are taking several days are mostly partially splitted over a few days. However, even the comparison is quite complicated, because the metrics are biased to the dataset, therefore some of the metrics are only comparable on the same dataset and are not transferable. This is why I mostly used CelebA. Also a comparison of higher resolutions might be a future thing here due to the exponential growing computational cost. Whereas it might be pretty interesting to observe the quality and impact of all modules on higher resolutions. But this brings us to the next project (future work diffusion models, to overcome a low resolution, but with the cost of extremely more computational power) Diffusion models. Up to now everything here might be obsolete, because it is an ongoing project, so don't expect much of it.

8 Training

This chapter is giving an introduction into the training and how the training was done. For each configuration a plot with three rows is represented. First row is showing the discriminator and generator loss. The following three are metrics, which approximate the performance of the training and give insights into the ability of creating undistinguishing data. The precision and recall from the best training iteration is shown later or in the appendix.




Figure 1: Resnet spectral normalization and feature matching




Figure 2: classical DCGAN




Figure 3: Resnet spectral normalization and hinge

9 Summary

10 Appendix